

# The Effect of Reward Function Design in RL-Based BTC Grid Trading: Pareto Frontier, Winner Reversal, and Out-of-Distribution Robustness of Prospect-Theoretic Reward

Chanhee Lee

Department of Computer Science

Seoul National University of Science and Technology

`happilyfly@seoultech.ac.kr`

June 2026

## Abstract

We quantitatively study the effect of *reward function design* on PPO-based reinforcement-learning policies for ATR-proportional grid trading on the BTC/USDT 1-hour market. Four reward variants—symmetric (**sym**), asymmetric ( $\beta = 3.42$ , **asym**), Differential Sharpe Ratio ( $\eta = 1/28$  h, **dsr**; [Moody and Saffell 2001](#)), and prospect-theoretic ( $\alpha = 0.68$ ,  $\lambda = 3.30$ , **pt**; [Kahneman and Tversky 1979](#))—are each trained with  $10 \text{ seeds} \times 1\text{M}$  steps and evaluated on three environments: (i) single-split validation (Val 2021–2023), (ii) 6-fold combinatorial purged cross-validation [[López de Prado, 2018](#)] 15 paths, and (iii) the unsealed out-of-sample Test set (2024–2026, BTC \$42K  $\rightarrow$  \$76K bull regime).

We report five key findings. **(1)** All four variants exceed the ATR baseline by +11–24% on Val, yet partition into two clusters  $\{\text{sym}, \text{dsr}\}$  (aggressive) vs.  $\{\text{asym}, \text{pt}\}$  (conservative); within-cluster Cohen’s  $|d| < 0.3$  versus across-cluster  $|d| > 0.8$ . **(2)** Mechanism: reward loss asymmetry directly determines trade frequency and holding time (policy distance ratio 2.22 $\times$ ), and the sliding-window memory structure of DSR results in 2–6 $\times$  higher hold rate than other variants. **(3)** Winner reverses across evaluation environments: Val=**sym**(1.87)  $\rightarrow$  CPCV=**dsr**(1.41,  $p < 0.001$ )  $\rightarrow$  Test=**pt**(0.34–0.37, consistent across two model sources,  $p < 0.002$ ). **(4)** The OOS robustness of **pt** arises from a short-holding mechanism: loss aversion  $\lambda = 3.30$  combined with concave gain  $\alpha = 0.68$  trains policies to exit within mean 1.4 h (max 6 h), thereby avoiding the sell-side timing risk of the unseen bull regime. In contrast, **dsr** learns long holding (mean 4.58 h, max 169 h = 7 days) and records the worst OOS performance—demonstrating that the same reward formulation underlies both in-sample advantage and OOS failure as two sides of the same coin. **(5)** The Val–Test distribution shift is highly significant (KS  $p < 10^{-10}$ , –27% variance), while policies’ actions remain nearly invariant between Val and Test ( $\Delta \leq 5\%$ ), indicating that the regime shift, not the policy, drives OOS differences.

Our contribution is twofold. First, we demonstrate quantitatively that reward design

exerts its effect through a *trade-off dimension in risk profile* rather than a single Sharpe “alpha source.” Second, we provide the first quantitative evidence that Kahneman-Tversky prospect theory confers out-of-distribution safety on RL trading policies. Code and data: <https://github.com/cosmicpotato2047/capstone-rl-trading>.

**Keywords:** Reinforcement Learning, Grid Trading, PPO, Reward Design, Prospect Theory, Differential Sharpe Ratio, Combinatorial Purged CV, Out-of-Distribution Generalization, Cryptocurrency

# Contents

|          |                                                                   |           |
|----------|-------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                               | <b>5</b>  |
| 1.1      | Motivation and Research Question . . . . .                        | 5         |
| 1.2      | Contributions . . . . .                                           | 5         |
| 1.3      | Paper Structure . . . . .                                         | 6         |
| <b>2</b> | <b>Related Work</b>                                               | <b>6</b>  |
| 2.1      | Market Making and Grid Trading . . . . .                          | 6         |
| 2.2      | Reinforcement Learning for Trading . . . . .                      | 7         |
| 2.3      | Reward-Design Theory . . . . .                                    | 7         |
| 2.4      | Evaluation Methodology . . . . .                                  | 8         |
| <b>3</b> | <b>Method</b>                                                     | <b>9</b>  |
| 3.1      | MDP Formulation . . . . .                                         | 9         |
| 3.2      | Trading-Environment Dynamics . . . . .                            | 10        |
| 3.3      | ATR Baseline . . . . .                                            | 11        |
| 3.4      | Four Reward Variants . . . . .                                    | 12        |
| 3.5      | PPO Training and the Stabilization Package . . . . .              | 13        |
| 3.6      | Evaluation Protocol . . . . .                                     | 13        |
| <b>4</b> | <b>Phase 1: Policy Saturation under Symmetric Reward</b>          | <b>14</b> |
| 4.1      | The Saturation Phenomenon . . . . .                               | 14        |
| 4.2      | Decisive Ablation: “RL = Fixed [1.0, 0.0]” . . . . .              | 15        |
| 4.3      | Mechanism: The ATR Formula Absorbs Degrees of Freedom . . . . .   | 15        |
| 4.4      | Motivation for This Paper . . . . .                               | 16        |
| <b>5</b> | <b>Pareto Frontier and Scenario D</b>                             | <b>16</b> |
| 5.1      | Experimental Setup . . . . .                                      | 16        |
| 5.2      | Per-Variant Results: Four Different Metric Winners . . . . .      | 17        |
| 5.3      | Pareto Frontier Visualization . . . . .                           | 17        |
| 5.4      | Statistical Separation of the Two Clusters . . . . .              | 17        |
| 5.5      | Policy Distance: Confirmation at the Action Level . . . . .       | 18        |
| 5.6      | Confirming Scenario D and Hypothesis Evaluation . . . . .         | 19        |
| 5.7      | Environment Dependence of the exp027_rl Prior Evidence . . . . .  | 20        |
| <b>6</b> | <b>Mechanism Quantification</b>                                   | <b>21</b> |
| 6.1      | Trajectory Dataset . . . . .                                      | 21        |
| 6.2      | Trade Frequency: Direct Effect of Reward Asymmetry . . . . .      | 21        |
| 6.3      | Hold Rate: The Mechanism of DSR’s Sliding-Window Memory . . . . . | 22        |
| 6.4      | Per-Regime Action Distribution . . . . .                          | 22        |
| 6.5      | Counterfactual State-grid Verification . . . . .                  | 23        |
| 6.6      | Summary of the Causal Chain . . . . .                             | 24        |

|          |                                                                     |           |
|----------|---------------------------------------------------------------------|-----------|
| <b>7</b> | <b>Robustness Analysis</b>                                          | <b>24</b> |
| 7.1      | Slippage Robustness . . . . .                                       | 24        |
| 7.1.1    | Setup . . . . .                                                     | 25        |
| 7.1.2    | Per-variant Results: Uniform $\sim 12\%$ Decay . . . . .            | 25        |
| 7.1.3    | MDD Is Almost Unchanged . . . . .                                   | 25        |
| 7.1.4    | Cluster Preservation: Scenario D Is Robust . . . . .                | 26        |
| 7.1.5    | Fair Comparison Against the ATR Baseline (Phase 16a) . . . . .      | 26        |
| 7.1.6    | Section Summary and Partial Verdict on H4 . . . . .                 | 26        |
| 7.2      | Multi-Split Evaluation via CPCV . . . . .                           | 27        |
| 7.3      | Out-of-Sample Evaluation and Prospect Theory's Robustness . . . . . | 27        |
| <b>8</b> | <b>Discussion</b>                                                   | <b>27</b> |
| <b>9</b> | <b>Conclusion</b>                                                   | <b>27</b> |

# 1 Introduction

Cryptocurrency markets exhibit micro-oscillations in hourly price that are largely orthogonal to long-term trends, motivating *grid trading* strategies that harvest revenue from volatility itself rather than from directional prediction. Grid trading is a simplified form of market making [Avellaneda and Stoikov, 2008], ranging from fixed-grid rules to volatility-adaptive variants that scale grid width by the Average True Range (ATR) [Wilder, 1978]. This work studies the latter: an ATR-proportional grid system whose grid width and limit-order placement are determined by a reinforcement-learning policy.

## 1.1 Motivation and Research Question

An earlier phase of this project hypothesized that PPO RL would discover dynamic volatility adaptation and outperform a fixed-grid baseline. Post-hoc behavioral analysis, however, revealed that the trained policy converged to a near-constant action [1.0, 0.0] (§4), and that the Bayesian-optimized ATR coefficient itself was the principal source of performance. This negative finding suggested that the symmetric-reward + ATR-proportional combination structurally absorbs the policy’s degrees of freedom, motivating the new research question of this paper:

**In BTC grid trading, how does the design of the reward function affect the behavioral pattern and generalization performance of RL policies? In particular, do loss-asymmetric (asymmetric, prospect-theoretic) or risk-adjusted-return (DSR) reward formulations yield meaningful differences in single-metric advantage or out-of-sample robustness?**

## 1.2 Contributions

This work makes the following four contributions.

1. **A statistically rigorous comparison framework for reward variants.** Four reward functions (`sym`, `asym`, `dsr`, `pt`) are trained in an identical environment with  $10 \text{ seeds} \times 1\text{M}$  steps, and pairwise statistical tests (Cohen’s  $d$ , bootstrap  $P(A > B)$ , IQM, 5% CVaR) are applied following Henderson et al. [2018]’s reproducibility guidelines.
2. **Scenario D — Pareto-frontier discovery.** Rather than a single “best variant” winner, the four variants statistically partition into two clusters  $\{\text{sym}, \text{dsr}\}$  vs.  $\{\text{asym}, \text{pt}\}$  forming a Pareto-like frontier on the Sharpe–MDD plane (§5). This outcome falls into none of the pre-registered scenario branches (A optimistic / B neutral / C pessimistic).
3. **Causal chain: reward  $\rightarrow$  behavior  $\rightarrow$  outcome.** Policy-level differences across variants are quantified on 1.04M step trajectories via a policy-distance matrix (within-cluster 0.129 vs. across-cluster 0.286), regime- specific trade/hold statistics, and action distribution comparison, confirming the mechanism reward loss asymmetry  $\rightarrow$  reduced trade frequency  $\rightarrow$  conservative cluster (§6).
4. **Winner reversal across evaluation environments and OOS robustness of prospect theory.** The winner reverses completely across single-split Val (`sym` 1<sup>st</sup>)  $\rightarrow$  CPCV multi-split

(**dsr** 1<sup>st</sup>)  $\rightarrow$  the unsealed Test (**pt** 1<sup>st</sup>, consistent across two sources  $p < 0.002$ ) (§7). The OOS robustness of **pt** is mechanistically attributed via trajectory analysis to its loss aversion  $\lambda = 3.30$  and concave utility  $\alpha = 0.68$  training policies to exit within mean 1.4h, thereby avoiding the sell-side timing risk of the unseen bull regime. This is, to our knowledge, the first quantitative demonstration that the Kahneman-Tversky prospect theory [Kahneman and Tversky, 1979] confers OOS safety on RL trading policies.

### 1.3 Paper Structure

Section 2 surveys related work on RL trading, reward design theory, and evaluation methodology. Section 3 details the environment design (MDP, ATR formula, four reward variants) and PPO training. Section 4 briefly reproduces the Phase 1 negative finding (“RL = Fixed [1.0, 0.0]”) as motivation. Section 5 reports the two-cluster Pareto-frontier result on Val (Scenario D), while Section 6 quantifies the mechanism. Section 7 presents the robustness trio: slippage (§7.1), CPCV multi-split (§7.2), and the central contribution of this paper—the Test OOS evaluation and the mechanistic basis of prospect theory’s robustness (§7.3). Section 8 discusses distribution shift quantification and limitations, and Section 9 concludes.

## 2 Related Work

This paper sits at the intersection of four areas: (i) market making and grid trading, (ii) reinforcement-learning-based trading, (iii) reward-design theory, and (iv) evaluation methodology for trading policies. This section surveys the key prior work in each area and clarifies where our contribution lies.

### 2.1 Market Making and Grid Trading

The academic ancestor of grid trading is the market-making model of Avellaneda and Stoikov [2008]. Their analysis derives a closed-form optimal policy for a risk-averse market maker who faces inventory risk and must place *asymmetric* bid/ask quotes around the mid-price. Two key insights are that (a) the bid–ask spread is a function of volatility and remaining time horizon, and (b) inventory skew shifts the quote midpoint asymmetrically. Grid trading is a simplified practical variant of this model in which orders are placed on a pre-defined ladder of price levels in order to monetize micro-volatility.

The standard mechanism for adapting grid widths to market volatility is the Average True Range [Wilder, 1978] (ATR; 14-period moving average is the default). ATR is a smoothed average of the true price range and provides a simple, robust rule that automatically widens the grid in volatility-cluster regimes. ATR-proportional grids of the form  $\Delta = c \cdot \text{ATR}$  (with proportional coefficient  $c$ ) are widely used in practice and serve as the baseline of this paper.

Most prior academic work in this area is limited to the analysis of *rule-based* policies, with relatively few attempts to learn the grid-width decision from data. This paper adopts a structure in which the proportional coefficient and the number of grid splits in the ATR formula are dynamically determined by an RL policy, while quantitatively analyzing how the design of

the reward function influences the resulting learned policy’s behavior—a combination, to our knowledge, that has not been explored.

## 2.2 Reinforcement Learning for Trading

The standard DRL baseline for cryptocurrency trading is the PPO/A2C/DQN comparison framework of Zhang et al. [2020], which feeds price time series directly into a PPO policy and shows consistent improvements over simple buy-and-hold (B&H) and moving-average crossover strategies. The ACM survey of Sun et al. [2023] and the more recent review of Liu [2025] both identify two open problems in the area: (a) *interpretability of policy behavior* and (b) *generalization across market regimes*. This paper addresses both directly: §6 quantifies the causal chain from reward form to behavioral difference, and §7 documents out-of-sample (OOS) generalization phenomena that are not captured by single-split validation alone.

There are four directly preceding works. Liu et al. [2020] feeds the output of an LSTM price predictor into the state of a PPO trading policy in a two-stage pipeline. The limitation of this approach is that prediction errors propagate directly into trading decisions; this paper bypasses the limitation by *eliminating* the price-prediction stage and using a volatility-adaptive grid structure instead. Yasin and Gill [2024] compares DQN/PPO/A2C with 20 technical indicators, but the action space is discrete buy/sell with no hold action. This paper adopts a continuous two-dimensional action (aggressiveness and profit target) that supports fine-grained grid control. Pham et al. [2025] uses DQN to select one of five predefined strategies as a meta-strategy, but the strategies themselves are not learned. Bandarupalli [2025] augments PPO with a volatility penalty and a transaction-cost term, but stays at the level of hyperparameter tuning within a single reward variant and does not compare across reward *forms*, which is the focus of this paper.

A particularly relevant prior work is Gort et al. [2022], which empirically shows that *single-split train/test results do not guarantee OOS generalization* in cryptocurrency DRL trading, and recommends CPCV (Combinatorial Purged Cross-Validation) and DSR (Deflated Sharpe Ratio). This paper adopts those recommendations, using Val (2021–2023), CPCV (6-fold inside the Train partition), and Test (2024–unsealed) environments concurrently. Our *winner reversal* finding (Val `sym`  $\rightarrow$  CPCV `dsr`  $\rightarrow$  Test `pt`) is a quantitative confirmation that Gort et al.’s concern reappears in the dimension of reward-variant selection.

In one sentence, our position is: whereas prior DRL trading literature focuses on comparing algorithms and state designs under a *single* reward function, this paper performs a controlled comparison of *four reward functions* on top of the same algorithm (PPO) and the same environment, and separately measures their effects on in-sample / multi-split / OOS environments.

## 2.3 Reward-Design Theory

The four reward variants compared in this paper draw from three theoretical lineages in reinforcement learning and behavioral economics. (i) Symmetric (`sym`) is the standard PnL-return reward and serves as our baseline. (ii) Asymmetric (`asym`) places a positive weight  $\beta > 1$  on losses, a simple first-order approximation of a risk-averse utility. (iii) Differential Sharpe Ratio

(**dsr**) is the formulation of [Moody and Saffell \[2001\]](#). (iv) Prospect-theoretic reward (**pt**) directly uses the utility function of [Kahneman and Tversky \[1979\]](#) and [Tversky and Kahneman \[1992\]](#),  $v(x) = \text{sgn}(x)|x|^\alpha$  for  $x \geq 0$  and  $-\lambda|x|^\alpha$  for  $x < 0$ , as a per-step reward.

The DSR formulation of [Moody and Saffell \[2001\]](#) uses the *differential change of the online Sharpe-ratio estimator* as the step reward, tracking first and second moments via exponentially-weighted moving averages (EWMA, decay rate  $\eta$ ). This has two implications for RL training: (a) it directly optimizes a risk-adjusted return rather than raw PnL, and (b) the sliding-window memory imposes a *multi-step temporal dependency* on the policy. As we show in §6, this memory structure is what causes the **dsr** policy to learn 2–6× longer holding times than the other variants—a property that becomes the in-sample advantage of §7 and the OOS failure of §7.3 as two sides of the same causal coin.

Prospect theory [[Kahneman and Tversky, 1979](#)] is a foundational model of human decision making under risk, with two characteristic parameters: (a) loss aversion  $\lambda \approx 2.25$  (standard value from [Tversky and Kahneman 1992](#)) and (b) the concavity exponent  $\alpha \approx 0.88$ . While the theory has been widely used to describe human trader behavior (e.g., the disposition effect of clinging to losses and realizing gains too early), *direct use of prospect-theoretic utility as an RL reward* is rare in the literature. To our knowledge, this paper is the first to directly implement the [Kahneman and Tversky \[1979\]](#) utility as an RL reward (with  $\alpha = 0.68$ ,  $\lambda = 3.30$ , tuned by Optuna) and quantitatively establish the OOS robustness that the resulting policy exhibits.

[Ng et al. \[1999\]](#)’s policy-invariance theorem of potential-based reward shaping states that “a monotone transformation of the reward signal does not alter the policy learned by RL.” At first glance, this paper’s comparison of four reward variants may appear to conflict with the theorem, but it does not: the **sym**  $\rightarrow$  **asym**/**pt** transformations are *non-potential-based* (they apply an asymmetric scale on the negative half-line), while **dsr** is not a simple per-step reward transformation but a *state-extended* reward (the sliding-window moments effectively become part of the state). Our reward-variant comparison thus lies outside the scope of the Ng et al. theorem, and indeed we confirm in §6 that policy behavior differs across variants in a statistically significant manner.

## 2.4 Evaluation Methodology

The two principal pitfalls in evaluating RL trading policies are (a) backtest overfitting and (b) seed dependence of results. [Henderson et al. \[2018\]](#) empirically demonstrates that seed variability alone can produce large differences for the same algorithm, and proposes 5–10 seed statistical comparisons, IQM (Interquartile Mean), and bootstrap confidence intervals as standard recommendations. This paper uses 10 seeds per variant (Val/exp033) and 100 models (Test), and applies pairwise Cohen’s  $d$  and bootstrap  $P(A > B)$  to all comparisons.

[López de Prado \[2018\]](#) (Ch. 8) proposes *Combinatorial Purged Cross-Validation* (CPCV), which avoids train/test leakage in time-series data while leveraging  $\binom{N}{k}$  split combinations to strengthen the generalization guarantees of single-split results. We adopt 6-fold CPCV (15 paths) in §7.2.

Bailey and López de Prado [2014]’s *Deflated Sharpe Ratio* ( $\text{DSR}^*$ ) is a statistical-significance test on the Sharpe ratio that corrects for the number of backtest trials and for the skewness/kurtosis of the underlying distribution. We apply Bonferroni 4-way correction together with  $\text{DSR}^*$  testing to the multiple comparisons of CPCV results (§7.2), and confirm that all four reward variants are significant relative to the ATR baseline at  $p < 0.001$ .

The methodological distinctiveness of this paper is the *simultaneous use of three evaluation environments* (Val / CPCV / Test) that allowed us to discover the winner reversal. Studies relying solely on single-split or solely on CPCV cannot expose OOS generalization failure, and studies relying solely on OOS cannot expose the in-sample conservatism of **pt**. The simultaneous application of three environments is, in our view, a methodological contribution: it generalizes the multi-evaluation recommendations of Henderson et al. [2018] and Gort et al. [2022] to the single dimension of reward-variant comparison.

## 3 Method

### 3.1 MDP Formulation

We cast single-asset dynamic grid trading on the BTC/USDT 1-hour series as a finite-horizon Markov decision process  $\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ . An episode runs for 720 hours (30 days) starting from a randomly chosen step within the training/validation range, and each step  $t$  corresponds to a 1-hour price bar.

**State**  $\mathcal{S} \subset \mathbb{R}^7$ . The observation  $\mathbf{s}_t = (s_t^{(0)}, \dots, s_t^{(6)})$  couples market and portfolio state:

$$s_t^{(0)} = \log(p_t / \bar{p}_{t,168}) \quad (\text{price vs. 1-week mean}) \quad (1)$$

$$s_t^{(1)} = \frac{\bar{p}_t^{\text{avg}} - p_t}{\bar{p}_t^{\text{avg}}} \quad (\text{divergence from average cost}) \quad (2)$$

$$s_t^{(2)} = \frac{h_t \cdot p_t}{C_0} \quad (\text{holdings value / capital}) \quad (3)$$

$$s_t^{(3)} = \frac{c_t}{C_0} \quad (\text{cash ratio}) \quad (4)$$

$$s_t^{(4)} = \frac{\text{ATR}_{168}(p)}{p_t} \quad (\text{volatility}) \quad (5)$$

$$s_t^{(5)} = \frac{p_t - p_{t-72}}{p_{t-72}} \quad (\text{short-term trend, 3d}) \quad (6)$$

$$s_t^{(6)} = \frac{p_t - p_{t-720}}{p_{t-720}} \quad (\text{long-term trend, 30d}) \quad (7)$$

Here  $p_t$  is the bar close,  $\bar{p}_t^{\text{avg}}$  is the average cost of the current position (set to 0 when no position is held, in which case  $s_t^{(1)}$  is also 0),  $h_t$  is the BTC quantity,  $c_t$  the cash balance, and  $C_0 = 10,000$  USDT is the initial capital.  $\text{ATR}_{168}$  is the Average True Range [Wilder, 1978] computed over a 168-bar (1-week) window, and  $\bar{p}_{t,168}$  is the corresponding simple moving average. All components are pre-normalized by a 168-bar rolling z-score computed on the training set and clipped to

$[-5, 5]$ .

**Action**  $\mathcal{A} = [0, 1]^2$ . At each step the policy emits a continuous 2D action  $\mathbf{a}_t = (a_t^{(0)}, a_t^{(1)})$ . The first coordinate  $a_t^{(0)}$  (*aggressiveness*) controls two buy limit orders, and  $a_t^{(1)}$  (*profit target*) controls two sell limit orders. We avoid a discrete buy/sell/hold action space because fine-grained grid width adjustment and volatility adaptation are most naturally expressed in continuous values; discretization would artificially limit the expressiveness of the action space.

**ATR-proportional grid formula.** Grid widths and positions are determined through volatility-normalized action coordinates. With  $r_t = \text{ATR}_{168}(p)/p_t$ , current price  $p_t$ , and average cost  $\bar{p}_t^{\text{avg}}$ ,

$$g_t^{\text{buy\_hi}} = r_t \cdot (A_b + B_b \cdot a_t^{(0)}) \quad (8)$$

$$g_t^{\text{buy\_lo}} = r_t \cdot (C_b + D_b \cdot a_t^{(0)}) \quad (9)$$

$$g_t^{\text{sell\_mkt}} = r_t \cdot (A_s + B_s \cdot a_t^{(1)}) \quad (10)$$

$$g_t^{\text{sell\_cost}} = r_t \cdot (C_s + D_s \cdot a_t^{(1)}) \quad (11)$$

Quote prices:

$$q_t^{\text{buy\_hi}} = p_t(1 - g_t^{\text{buy\_hi}}), \quad q_t^{\text{buy\_lo}} = p_t(1 - g_t^{\text{buy\_lo}}) \quad (12)$$

$$q_t^{\text{sell\_mkt}} = p_t(1 + g_t^{\text{sell\_mkt}}), \quad q_t^{\text{sell\_cost}} = \bar{p}_t^{\text{avg}}(1 + g_t^{\text{sell\_cost}}) \quad (13)$$

The coefficients  $\{A_b, B_b, C_b, D_b, A_s, B_s, C_s, D_s\}$  are fixed environment constants, with values reported in §3.3. Only the `sell_cost` quote is anchored at the average cost; the choice explicitly exposes “break-even recovery” to the policy.

**Reward.** The baseline (`sym`, symmetric) form is the 1-step PnL ratio:

$$r_t^{\text{sym}} = \frac{E_t - E_{t-1}}{C_0}, \quad E_t = c_t + h_t \cdot p_t \quad (14)$$

Transaction fees are already deducted from  $c_t$  inside `_execute_buy` and `_execute_sell`, so no separate subtraction appears in the reward. All four reward variants in this paper are defined as functions of  $r_t^{\text{sym}}$  and detailed in §3.4.

## 3.2 Trading-Environment Dynamics

This subsection details order execution, the cycle definition, and the transaction-cost model.

**Limit-order execution.** Fill decisions for the four quotes use the next bar’s high/low ( $p_{t+1}^H, p_{t+1}^L$ ):

$$\begin{aligned} p_{t+1}^L &\leq q_t^{\text{buy\_hi}} \Rightarrow \text{fill at } q_t^{\text{buy\_hi}} \\ p_{t+1}^L &\leq q_t^{\text{buy\_lo}} \Rightarrow \text{fill at } q_t^{\text{buy\_lo}} \\ p_{t+1}^H &\geq q_t^{\text{sell\_mkt}} \Rightarrow \text{fill at } q_t^{\text{sell\_mkt}} \\ p_{t+1}^H &\geq q_t^{\text{sell\_cost}} \Rightarrow \text{fill at } q_t^{\text{sell\_cost}} \end{aligned} \quad (15)$$

When both buy and sell quotes are filled in the same bar we apply a *sell-first* rule, prioritizing inventory liquidation. The fill price is the *quote price itself*, not the next-bar extreme. Filling at the next-bar extreme injects favorable bias into the simulation; we therefore discarded an earlier “next\_low/next\_high” execution rule and use only the limit-order rule throughout this paper.

**Cycle definition.** A cycle starts when the *first buy fills while holdings = 0* and ends when subsequent sells return holdings to 0. Multiple partial buys/sells can occur inside a single cycle, and the cycle return is  $\text{cycle\_pnl} = (c_{\text{end}} - c_{\text{start}})/c_{\text{start}}$ . At the start of each cycle, the slot size used for the cycle is set:

$$\text{slot\_size}_{\text{cycle}} = \frac{c_{\text{start}}}{n_{\text{splits}}}, \quad \text{order\_size} = \frac{\text{slot\_size}}{n_{\text{buy\_orders}}} \quad (16)$$

We use  $n_{\text{splits}} = 2$  and  $n_{\text{buy\_orders}} = 2$  (the best values from Phase 2 Optuna Trial #34). `threshold_basis` is fixed to `avg_price`, meaning the “liquidate all” detection uses the average cost.

**Transaction-cost model.** Transaction costs are set to *0.05% per side* (Binance maker fee). On a buy, the filled quantity is  $\Delta h = \text{order\_size} \cdot (1 - f)/q$  and the cash change is  $\Delta c = -\text{order\_size}$ , so the fee is already reflected in  $c_t$ . In §7.1 (exp033) we add an additional slippage  $\sigma = 0.02\%$  that modifies fill prices to  $q^{\text{buy}} \cdot (1 + \sigma)$  and  $q^{\text{sell}} \cdot (1 - \sigma)$ . Slippage is not applied to the base environment and appears only in the robustness analysis.

### 3.3 ATR Baseline

The comparison baseline *shares the environment and grid formula* with our RL policies but determines actions via a rule. Concretely, the ATR baseline reuses the grid formulas of §3.1 for  $g^{\text{buy\_hi}}, g^{\text{buy\_lo}}, g^{\text{sell\_mkt}}, g^{\text{sell\_cost}}$  with all RL action degrees of freedom *zeroed* ( $a_t^{(0)} = a_t^{(1)} = 0$ ), so that the  $B_b, D_b, B_s, D_s$  terms drop and the grid widths are fixed time-independent constants determined entirely by  $A_b, C_b, A_s, C_s$ . The coefficient values come from Phase 2 Optuna Trial #34:

Table 1: Grid coefficients of the ATR baseline (Phase 2, Optuna Trial #34).

| Coef. | Value  | Role                      | Coef. | Value | Role                   |
|-------|--------|---------------------------|-------|-------|------------------------|
| $A_b$ | 1.665  | buy_hi base width         | $A_s$ | 0.285 | sell_mkt base width    |
| $C_b$ | 6.070  | buy_lo base width         | $C_s$ | 1.951 | sell_cost base width   |
| $B_b$ | 1.748  | buy_hi action sensitivity | $B_s$ | 1.950 | sell_mkt action sens.  |
| $D_b$ | 18.683 | buy_lo action sensitivity | $D_s$ | 7.500 | sell_cost action sens. |

The ATR baseline yields Val Sharpe 1.378 without slippage,<sup>1</sup> dropping to 0.835 when slippage of 0.02% is added.

<sup>1</sup>In the early RESEARCH\_LOG (May 14, 2026) we recorded 1.505; the value was re-measured to 1.378 after the Phase 16a unification of the evaluation setup (`nval_episodes`, `env config synchronization`). We use 1.378 throughout the paper.

### 3.4 Four Reward Variants

We now define the four reward variants that form the principal comparison unit of this paper. All are expressed as functions of  $x_t := r_t^{\text{sym}} = (E_t - E_{t-1})/C_0$  and differ *only* in functional form; the state, action, ATR grid, and fee model are otherwise identical.

(i) **Symmetric (sym).**

$$r_t^{\text{sym}} = x_t = \frac{E_t - E_{t-1}}{C_0} \quad (17)$$

No hyperparameter. Baseline of this paper.

(ii) **Asymmetric (asym).**

$$r_t^{\text{asym}} = \begin{cases} x_t & \text{if } x_t \geq 0 \\ \beta \cdot x_t & \text{if } x_t < 0 \end{cases}, \quad \beta > 1 \quad (18)$$

A first-order approximation of a risk-averse utility, weighting losses by  $\beta$ . An Optuna 30-trial  $\times$  200k-step search over the range  $[1.0, 4.0]$  to maximize Val Sharpe selected the best  $\beta = 3.420$ .

(iii) **Differential Sharpe Ratio (dsr).** The formulation of [Moody and Saffell \[2001\]](#). Given an EWMA decay rate  $\eta$ , online estimates of the first and second return moments evolve as

$$A_t = A_{t-1} + \eta(x_t - A_{t-1}) \quad (19)$$

$$B_t = B_{t-1} + \eta(x_t^2 - B_{t-1}) \quad (20)$$

$$\Delta A_t = x_t - A_{t-1}, \quad \Delta B_t = x_t^2 - B_{t-1} \quad (21)$$

The DSR step reward is the differential of the on-line Sharpe estimator  $\hat{S}_t = A_t/\sqrt{B_t - A_t^2}$ :

$$r_t^{\text{dsr}} = \frac{B_{t-1} \Delta A_t - \frac{1}{2} A_{t-1} \Delta B_t}{(B_{t-1} - A_{t-1}^2)^{3/2}} \quad (22)$$

A 30-trial Optuna search over  $[1/720, 1/24]$  (log-uniform) selected the best  $\eta = 0.0352 \approx 1/28$  h, an EWMA window of approximately 1.2 days.

(iv) **Prospect-theoretic (pt).** The utility function of [Kahneman and Tversky \[1979\]](#) and [Tversky and Kahneman \[1992\]](#) applied directly as a step reward:

$$r_t^{\text{pt}} = \begin{cases} \text{sgn}(x_t) |x_t|^\alpha & \text{if } x_t \geq 0 \\ -\lambda |x_t|^\alpha & \text{if } x_t < 0 \end{cases}, \quad 0 < \alpha \leq 1, \quad \lambda \geq 1 \quad (23)$$

$\alpha < 1$  encodes the concavity of the utility (diminishing marginal utility of gains), and  $\lambda > 1$  encodes loss aversion. The standard values from [Tversky and Kahneman \[1992\]](#), fitted to human behavior, are  $\alpha = 0.88$  and  $\lambda = 2.25$ . The Val-Sharpe-maximizing values for our RL policy, however, are  $\alpha = 0.683$  and  $\lambda = 3.303$ , found by Optuna 30 trials over  $\alpha \in [0.5, 1.0]$  and  $\lambda \in [1.0, 4.0]$ . The fact that RL prefers a *more concave* utility and *stronger loss aversion* than

the human standard is consistent with the mechanism in §6 and the OOS robustness reported in §7.3.

**Optuna search summary.** For each of the three variants `asym`, `dsr`, `pt` we ran  $30 \text{ trials} \times 200\text{k steps}$  using a TPE sampler [Akiba et al., 2019] (seed=42, MedianPruner with  $n_{\text{startup}} = 5$ ). The total training cost was  $3 \times 30 \times 200\text{k} = 18\text{M steps}$ , taking approximately 2 hours on a Windows 11 CPU with  $n_{\text{envs}} = 4$ . `sym` has no hyperparameter and skips this stage.

### 3.5 PPO Training and the Stabilization Package

Policy training uses PPO [Schulman et al., 2017] from stable-baselines3 [Raffin et al., 2021]. Standard PPO hyperparameters ( $\gamma = 0.99$ ,  $\lambda_{\text{GAE}} = 0.95$ , clip  $\epsilon = 0.2$ ,  $v_f$  coef = 0.5, gradient clip = 0.5, MLP policy [64, 64], batch=256,  $n_{\text{epochs}} = 10$ ,  $n_{\text{steps}} = 4096$ ,  $n_{\text{envs}} = 4$ ) follow stable-baselines3 defaults. To mitigate the reward sparsity of our environment (the per-step reward is near 0 when no fill occurs) and the late-stage policy collapse observed in Phase 2 (exp020–022), we add the following *stabilization package* (exp030):

1. **Learning-rate linear decay:**  $\alpha_t$  linearly decays from  $3 \times 10^{-4}$  to  $1 \times 10^{-5}$  over training. This bounds the size of late-stage policy updates and dampens post-best collapse.
2. **Entropy coefficient annealing:**  $c_{\text{ent}}$  linearly anneals from 0.01 to 0.001, ensuring sufficient initial exploration while encouraging convergence to a near-deterministic policy late in training.
3. **Target KL early stop:** `target_kl = 0.02` measures the KL divergence after each minibatch update and aborts the update for that batch if the threshold is exceeded, preventing PPO from making myopic large policy changes.
4. **Best-model checkpoint:** Val Sharpe is evaluated every 50k steps ( $n_{\text{eval\_episodes}} = 5$ ) and the weights at the best Val Sharpe are saved separately, providing protection against late-stage collapse.

The package was developed in exp030 (with `sym` reward) and produced Val Sharpe best 1.974 at step 550k, falling to final 1.209 at step 1M. Because the collapse pattern after step 700k persists despite the package, *all evaluations in this paper use the best checkpoint*; the final model is never used.

### 3.6 Evaluation Protocol

We use *three evaluation environments concurrently*. The data partition is strict time-order to prevent leakage:

Table 2: Data partition.

| Partition          | Range                   | Bars (approx.) |
|--------------------|-------------------------|----------------|
| Train              | 2017-08-17 – 2020-12-31 | ~30,000        |
| Val (single-split) | 2021-01-01 – 2023-12-31 | ~26,000        |
| Test (sealed)      | 2024-01-01 – 2026-04    | ~20,000        |

**Val (single-split).** The main comparison environment of §5. For each reward variant we train  $10 \text{ seeds} \times 1\text{M steps}$  and evaluate the `best_model` on the entire Val range with  $n_{\text{eval\_episodes}} = 5$ , producing per-seed Sharpe/MDD/Calmar/trade-count/return.

**CPCV 6-fold.** The multi-split environment of §7.2. Following López de Prado [2018], we apply 6-fold Combinatorial Purged Cross-Validation *inside the Train partition*, producing  $\binom{6}{2} = 15$  paths (each consisting of 5 months of training plus 1 month of evaluation). A 1-week purge gap between train and evaluation folds prevents information leakage. We execute  $4 \text{ variants} \times 15 \text{ paths} \times 1 \text{ seed} = 60$  runs.

**Test (unsealed).** The OOS environment of §7.3. For honesty, following Gort et al. [2022], Test was *kept sealed from any search or hyperparameter decision* until exp035 (May 16, 2026). After unsealing, we evaluate 100 RL models ( $4 \text{ variants} \times \text{multi-seed/seed-pair}$ ) along with the ATR baseline and Buy-and-Hold on Test.

**Statistical tests.** All pairwise comparisons report Cohen’s  $d$  and bootstrap  $P(A > B)$  ( $10^4$  resamples). Single-sided t-tests against ATR on CPCV use Bonferroni 4-way correction along with the Deflated Sharpe Ratio test of Bailey and López de Prado [2014]. To report evaluation variability and distributional skew, we follow Henderson et al. [2018] and add IQM (Interquartile Mean) and 5% CVaR as secondary metrics.

## 4 Phase 1: Policy Saturation under Symmetric Reward

This section briefly reproduces the negative finding of an *earlier* Phase 1 (April 2026) experiment series exp020–022 that precedes this paper. The finding is the starting point that motivates the reward-formulation comparison of §3.4. The Phase 1 experiments were run on *Env-v2* (2D ATR-proportional formula, next-bar-extreme execution) rather than on this paper’s canonical *Env-v4*, so the absolute Sharpe values include a favorable-bias artifact. However, the qualitative observation at the heart of this section—that the *policy converges to a constant*—is environment-independent and remains valid as motivation.

### 4.1 The Saturation Phenomenon

The Phase 1 setup is almost identical to §3 of this paper: state 7D, action 2D, ATR-proportional grid formula, symmetric reward  $r_t = (E_t - E_{t-1})/C_0$ . The two differences are (i) execution uses the next-bar extremes rather than limit fills, and (ii) the stabilization package of §3.5 had not yet been introduced.

Three experiments converged to the same conclusion.

**(a) exp020 (budget\_fraction).**  $a_t^{(0)}$  was redefined from “aggressiveness” to “fraction of cash budget”: at the start of each cycle, the RL policy decides how much of the cash balance to commit to the cycle ( $0 = \text{no entry}$ ,  $1 = \text{full commitment}$ ). The hypothesis was “commit less in

bear regimes and more in bull regimes,” but the learned policy converged to  $a_t^{(0)} \approx 1.0$  across all regimes.

**(b) exp021 (entry\_gate).**  $a_t^{(0)}$  was redefined as a “cycle entry gate” (0 = no entry, 1 = allow entry). The hypothesis was “block entries in bear regimes,” but the learned entry rate in the **bear** regime was 99.7%, effectively keeping the gate open at every bar.

**(c) exp022 (aggressiveness, identical semantics to this paper).**  $a_t^{(0)}$  was restored to “aggressiveness” and PPO was trained for 1M steps. Behavioral statistics of the deterministic policy: across all regimes  $a_t^{(0)} \approx 0$  and  $a_t^{(1)} \approx 0$  (both medians equal to 0.0000, means  $\leq 0.001$ ). The most decisive diagnostic is that the *raw network outputs* were  $[-9.19, -4.30]$ , deep in the saturation region of the squashing nonlinearity.

## 4.2 Decisive Ablation: “RL = Fixed [1.0, 0.0]”

We took the converged behavioral statistics and ran an ablation that replaces the RL policy with a *fixed* action. In the **aggressiveness** semantics,  $a_t = [1.0, 0.0]$  corresponds to the learned policy’s behavior ( $a_t^{(0)} \rightarrow 1$  produces the narrowest buy spacing in the ATR grid formula and  $a_t^{(1)} \rightarrow 0$  produces the narrowest sell spacing).

Table 3: Phase 1 (Env-v2) decisive ablation. Absolute Sharpe values reflect the next-bar-extreme execution favorable bias, but the *agreement* between RL and Fixed [1.0, 0.0] is environment-independent.

| Policy                                   | Val Sharpe    | Test Sharpe |
|------------------------------------------|---------------|-------------|
| RL exp020 ( <b>sym</b> reward, 1M steps) | <b>45.390</b> | 42.090      |
| Fixed [1.0, 0.0] (RL convergence point)  | <b>45.390</b> | 41.769      |
| Fixed [1.0, 0.5]                         | 4.116         | 4.683       |
| Fixed [1.0, 1.0]                         | 1.060         | 1.843       |

The decisive observation in Table 3 is that RL and Fixed [1.0, 0.0] agree on Val Sharpe *to three decimal places*, meaning the result of 1M steps of PPO training reduces to a constant policy. The remaining Fixed policies [1.0, 0.5] and [1.0, 1.0] are substantially worse, partly justifying the claim that “RL did pick a meaningful point in the action space.” The crucial qualifier is that the point is time-independent.

## 4.3 Mechanism: The ATR Formula Absorbs Degrees of Freedom

The structural cause of policy saturation lies in the grid formula itself. Recalling from §3.1, the sell-side grid width is

$$g_t^{\text{sell\_mkt}} = \frac{\text{ATR}_{168}(p)}{p_t} \cdot (A_s + B_s \cdot a_t^{(1)}) \quad (24)$$

and the term  $\text{ATR}_{168}/p_t$  *already* reflects market volatility. In a bear regime (ATR $\uparrow$ ) the grid widens automatically, producing a conservative position; in a bull regime (ATR $\downarrow$ ) the grid narrows, producing an aggressive position. The role left for RL is “how should I further adjust the width that ATR already chose?”—and under a symmetric reward the only optimization

direction is *maximize cycle count*  $\rightarrow$  *maximize compounding*, whose unique optimum is to pick the narrowest buy spacing ( $a_t^{(0)} \rightarrow 1$ ) and the narrowest sell spacing ( $a_t^{(1)} \rightarrow 0$ ). There is no incentive to learn a state-dependent action.

#### 4.4 Motivation for This Paper

Read literally, the Phase 1 negative finding suggests “RL adds no value to BTC grid trading.” This paper takes a different stance: that the conclusion is conditional on the specific assumption of a *symmetric reward*. That is:

*If the combination of symmetric reward and ATR-proportional formula absorbs the RL degrees of freedom, can reformulating the reward as asymmetric (asymmetric, prospect-theoretic) or path-dependent (sliding-window DSR) restore the value-add of RL?*

This is precisely the hypothesis formalized in this paper’s main comparison (§3.4, §5). We will further see in §5 that, on the canonical environment Env-v4 (limit-order execution plus the stabilization package), even the symmetric reward yields a policy that exceeds the ATR baseline (Val Sharpe sym = 1.87 > ATR 1.378). The strong Phase 1 negative “RL = Fixed” is therefore a joint result of favorable-bias execution and insufficient training stabilization; once the environment is normalized and training is stabilized, the reward form itself becomes the principal channel through which RL adds value.

### 5 Pareto Frontier and Scenario D

This section presents the first main finding of this paper: the *Pareto-like frontier structure of the four reward variants*. Training each of the four reward forms defined in §3.4 for 10 seeds  $\times$  1M steps on the canonical environment (Env-v4) places the 40 resulting policies in two clusters that form a trade-off frontier in the Sharpe–MDD plane rather than collapsing onto a single winner. None of the three pre-registered scenarios (A/B/C) matches this outcome, so we define an ex-post scenario “D” as the honest summary, and it forms the starting point for the multi-dimensional quantification of the rest of the paper.

#### 5.1 Experimental Setup

We train each of the four variants of §3.4 with 10 seeds (42–51). PPO settings and the stabilization package follow §3.5 verbatim; hyperparameters for `asym`, `dsr`, and `pt` are the Optuna bests from §3.4 ( $\beta = 3.42$ ,  $\eta = 0.0352$ ,  $\alpha = 0.683$ ,  $\lambda = 3.303$ ). `sym` has no hyperparameter. The total training cost is  $4 \times 10 \times 1\text{M} = 40\text{M}$  steps, completed in approximately 3 hours 44 minutes on a Windows 11 CPU ( $n_{\text{envs}} = 4$ ).

Evaluation uses the `best_checkpoint` for each seed on the single-split Val window (2021–2023) of §3.6. Metrics are best Val Sharpe, final Val Sharpe (Val Sharpe at the 1M step), MDD, Calmar, trade count, and cumulative return. The gap between best and final Val Sharpe is a proxy for late-stage policy stability.

## 5.2 Per-Variant Results: Four Different Metric Winners

Table 4 reports the 10-seed mean  $\pm$  standard deviation for the six metrics.

Table 4: exp032b results. 10 seeds  $\times$  1M steps, evaluated on Val. The ATR baseline is a single deterministic policy so it has no std. **Bold** indicates per-metric leaders.

| Variant      | Best Sharpe                        | Final Sharpe                       | MDD (%)                           | Calmar       | Trades     | Return (%)  |
|--------------|------------------------------------|------------------------------------|-----------------------------------|--------------|------------|-------------|
| <b>sym</b>   | <b><math>1.871 \pm 0.22</math></b> | $1.015 \pm 0.43$                   | $3.27 \pm 0.82$                   | 0.60         | <b>120</b> | 6.60        |
| <b>dsr</b>   | $1.809 \pm 0.21$                   | <b><math>1.204 \pm 0.41</math></b> | $4.33 \pm 1.70$                   | 0.47         | 117        | <b>7.40</b> |
| <b>asym</b>  | $1.681 \pm 0.10$                   | $1.101 \pm 0.27$                   | <b><math>2.28 \pm 0.31</math></b> | <b>0.755</b> | 96         | 5.23        |
| <b>pt</b>    | $1.667 \pm 0.09$                   | $1.082 \pm 0.18$                   | $2.31 \pm 0.29$                   | 0.735        | 85         | 4.88        |
| ATR baseline | 1.378                              | —                                  | 9.83                              | 0.153        | 2,121      | 35.80       |

Two key observations follow.

(i) **All four variants exceed the ATR baseline.** Best Val Sharpe values range from +21% (**sym**) to +21% (**pt**) relative to ATR (1.378), and all are statistically significant (per-variant bootstrap 95% CI lower bound exceeds 1.5). This contrasts qualitatively with the “RL = Fixed [1.0, 0.0]” result of §4, and we attribute the difference to environment normalization (limit-order execution) and the stabilization package (§3.5). The fact that even **sym** clearly exceeds ATR rejects **H1** (**sym** RL  $\approx$  ATR).

(ii) **Four different metric leaders.** **sym** leads on best Sharpe, **dsr** leads on final Sharpe and cumulative return, **asym** leads on MDD and Calmar, and **pt** sits a hair behind **asym** on all metrics. There is no single winner; the choice of metric determines which variant is “best.” This is a multi-dimensional trade-off structure, and **H2 strong** (**asym/dsr/pt** > **sym**) is **partially rejected**.

## 5.3 Pareto Frontier Visualization

Figure 1 plots the 40 policies’ (Sharpe, MDD) coordinates color-coded by reward variant. Two features stand out.

First, within each variant the 10-seed points are tightly grouped: **asym** and **pt** have std  $\approx$  0.10 (std/mean  $\approx$  6%), indicating strong seed robustness for the conservative variants; **sym** and **dsr** have std  $\approx$  0.22, slightly more spread.

Second, the point sets for **sym** and **dsr** *largely overlap*, and so do those for **asym** and **pt**, but the two pairs hardly overlap with each other. The visual point-level separation is quantified statistically in the next subsection.

## 5.4 Statistical Separation of the Two Clusters

Pairwise Cohen’s  $d$  on best Val Sharpe and the bootstrap  $P(A > B)$  ( $10^4$  resamples) appear in Tables 5 and 6.

The pattern in Table 5 is decisive.




Figure 1: Scatter of the 40 policies (4 variants  $\times$  10 seeds) in the Sharpe–MDD plane. The horizontal axis is best Val Sharpe (higher better) and the vertical axis is MDD (lower better). **sym**/**dsr** cluster in the *aggressive* region of high Sharpe + high MDD; **asym**/**pt** cluster in the *conservative* region of moderate Sharpe + low MDD. The ATR baseline sits at (1.378, 9.83) in the same plane and is Pareto-dominated by every RL policy.

Table 5: Pairwise Cohen’s  $d$  (best Sharpe, row–col).  $|d| < 0.3$  is a “small effect”;  $|d| > 0.8$  is a “large effect.” Within-cluster cells with  $|d| < 0.3$  and across-cluster cells with  $|d| > 0.8$  are bolded.

|      | sym          | asym         | dsr          | pt           |
|------|--------------|--------------|--------------|--------------|
| sym  | —            | <b>+1.10</b> | +0.29        | <b>+1.19</b> |
| asym | <b>−1.10</b> | —            | −0.79        | +0.15        |
| dsr  | −0.29        | +0.79        | —            | <b>+0.89</b> |
| pt   | <b>−1.19</b> | −0.15        | <b>−0.89</b> | —            |

- **Within cluster:** sym vs dsr  $|d| = 0.29$ , asym vs pt  $|d| = 0.15$ . Both are “small effects.”
- **Across cluster:** sym vs asym  $|d| = 1.10$ , sym vs pt  $|d| = 1.19$ , dsr vs asym  $|d| = 0.79$ , dsr vs pt  $|d| = 0.89$ . All four pairs have  $|d| > 0.79$ .

Hence the four variants statistically partition into:

$$\underbrace{\{\text{sym}, \text{dsr}\}}_{\text{aggressive}} \quad \text{vs.} \quad \underbrace{\{\text{asym}, \text{pt}\}}_{\text{conservative}}$$

## 5.5 Policy Distance: Confirmation at the Action Level

The statistics above describe separation at the level of Sharpe distributions. To check whether the clustering also exists at the *policy-action* level, we deterministically rolled out each of the 40

Table 6: Pairwise  $P(\text{row} > \text{col})$  on best Sharpe ( $10^4$  bootstrap resamples). Values near 1 indicate dominance by the row variant; 0.5 indicates equivalence.

|      | sym   | asym  | dsr   | pt    |
|------|-------|-------|-------|-------|
| sym  | —     | 0.998 | 0.746 | 0.999 |
| asym | 0.003 | —     | 0.031 | 0.634 |
| dsr  | 0.254 | 0.968 | —     | 0.984 |
| pt   | 0.001 | 0.363 | 0.018 | —     |

best\_models across the entire Val range, recorded their action time series, and computed the average  $L_2$  action distance between each variant pair (averaged over 10 seeds per variant).




Figure 2: Pairwise policy distance matrix for the four variants. Each cell is the mean  $L_2$  action distance between two variants’ policies. Diagonal entries are 0; darker shading indicates closer policies.

The off-diagonal values in Figure 2:

- **Within cluster:** sym vs dsr = 0.123, asym vs pt = 0.134. Mean = 0.129.
- **Across cluster:** sym vs asym = 0.209, sym vs pt = 0.326, dsr vs asym = 0.256, dsr vs pt = 0.353. Mean = 0.286.

The **across/within** ratio is  $0.286/0.129 = 2.22\times$ : policies within the same cluster are  $2.22\times$  closer to one another (in action space) than policies across clusters. The Sharpe-level cluster structure thus aligns with the behavior-level cluster structure.

## 5.6 Confirming Scenario D and Hypothesis Evaluation

To evaluate results without ex-post rationalization, we pre-registered the following scenario branches.

- **Scenario A (optimistic):** asym/dsr/pt exceed both sym and ATR with Cohen’s  $d \geq 0.5$ . Directly supports “reward design is a principal channel of RL alpha.”
- **Scenario B (neutral):** variants differ but none exceeds ATR. “The ATR-proportional formula absorbs volatility” — an extension of the negative finding.
- **Scenario C (pessimistic):** variants are indistinguishable. “Reward reformulation adds no alpha” — a strengthening of the Phase 1 finding.

Our results fit none of the pre-registered branches. All four variants exceed ATR (rejecting B), but there is no single winner (rejecting the optimistic A’s “asym/dsr/pt lead” form), and the variants are not indistinguishable (rejecting C, given the two statistically separated clusters). We therefore define an ex-post scenario **D (Pareto frontier / multi-dimensional trade-off)**:

*The four reward variants do not differ as a single-metric alpha source but as a **trade-off along risk-profile dimensions**, forming a Pareto-like frontier on the Sharpe–MDD plane. The two clusters ( $\{\text{sym}, \text{dsr}\}$  and  $\{\text{asym}, \text{pt}\}$ ) are consistently identified both at the level of Sharpe distributions (within-cluster  $|d| < 0.30$ , across-cluster  $|d| > 0.79$ ) and at the level of policy actions (distance ratio  $2.22\times$ ).*

Per-hypothesis verdicts at this point in the paper:

- **H1 (sym RL  $\approx$  ATR): rejected.** sym’s best Sharpe 1.871 exceeds ATR 1.378 with  $|d| > 0.6$  (per-variant t-test  $p < 10^{-3}$ ).
- **H2 weak (asym/dsr/pt  $>$  ATR): supported.** All four variants exceed ATR; the conservative cluster additionally wins MDD/Calmar.
- **H2 strong (asym/dsr/pt  $>$  sym): partially rejected.** dsr is statistically indistinguishable from sym (within cluster), and asym/pt have lower Sharpe than sym.
- **H3 (selective entry  $\rightarrow$  conservative): supported.** The conservative cluster has approximately 75% of sym’s trade count (asym 96, pt 85 vs. sym 120). The full mechanism is quantified in §6.

## 5.7 Environment Dependence of the exp027\_rl Prior Evidence

Part of the motivation for this paper was the strong asym ( $\beta = 2.0$ ) advantage observed in late Phase 2 (April 2026). On Env-v3 (4D absolute-gap, limit-order execution), the asym RL policy reached Test Sharpe 1.955, +109% above ATR Test 0.935. On our canonical Env-v4 (2D ATR-proportional, Optuna-retuned  $\beta = 3.42$ ), asym’s Val Sharpe is 1.681, *below* sym’s 1.871. The strong prior asym advantage thus *does not reproduce*. This indicates that the environment configuration (action dimensionality, grid formula form) has a larger effect than the reward variant. We honestly acknowledge this environment dependence in §8. The main conclusion of this section (Scenario D, the Pareto frontier) holds independently and is statistically established on canonical Env-v4.

## 6 Mechanism Quantification

§5 showed *what* happened: the four reward variants statistically separate into two clusters. This section quantifies *why*. The core hypothesis is that the loss asymmetry of the reward form directly determines the policy’s trade frequency and holding time, and that this behavioral difference is the causal origin of the risk-profile difference. We analyze the 1.04M-step trajectories of the 40 models to verify this causal chain step by step.

### 6.1 Trajectory Dataset

We replay the 40 best\_models of §5 on the full Val range in deterministic mode and record at every step the (state, action, reward, fill, holdings, regime label). Evaluation runs with `random_start=False` and without episode truncation, giving roughly 26,074 steps per seed in a single uninterrupted trajectory. The total is  $40 \times 26,074 \approx 1.04\text{M}$  rows, collected in 7.2 minutes. We run the following five analyses on this dataset: (1) Sharpe–MDD scatter [used in §5], (2) per-regime action distribution, (3) state-grid counterfactual, (4) per-regime behavior statistics (trade/hold rate), (5) policy distance [used in §5].

Volatility regimes are assigned by quantiles of  $s_t^{(4)} = \text{ATR}_{168}/p_t$ , partitioning the time axis into three bins (low/mid/high). Each (variant, regime) cell contains  $\sim 28,000$ – $30,000$  steps.

### 6.2 Trade Frequency: Direct Effect of Reward Asymmetry

Table 7: Per-regime trade/hold rate, averaged over 40 models (exp032c Menu 4). **Trade rate** = (steps with a fill) / (total steps). **Hold rate** = (steps with non-zero holdings and no fill) / (total steps). *Note the 2–6 $\times$  gap in dsr’s hold rate.*

| Variant | Trade rate |         |          | Hold rate    |              |              |
|---------|------------|---------|----------|--------------|--------------|--------------|
|         | low_vol    | mid_vol | high_vol | low_vol      | mid_vol      | high_vol     |
| sym     | 0.073      | 0.058   | 0.047    | 0.074        | 0.058        | 0.048        |
| dsr     | 0.073      | 0.057   | 0.047    | <b>0.142</b> | <b>0.120</b> | <b>0.120</b> |
| asym    | 0.058      | 0.044   | 0.038    | 0.036        | 0.028        | 0.023        |
| pt      | 0.049      | 0.037   | 0.032    | 0.031        | 0.023        | 0.020        |




Figure 3: Trade/hold rate over 4 variants  $\times$  3 vol regimes. Top row: trade rate. Bottom row: hold rate. `dsr`’s hold rate stands out at roughly 2–6 $\times$  the other variants.

Two key patterns emerge in Table 7 and Figure 3.

(i) **Trade-frequency ordering:  $\text{sym} \approx \text{dsr} > \text{asym} > \text{pt}$ .** In every regime,  $\{\text{sym}, \text{dsr}\}$  (the aggressive cluster) trade roughly 30–50% more often than  $\{\text{asym}, \text{pt}\}$  (the conservative cluster). The variant-induced effect (variant  $|d| \approx 1.0$ ) is *larger* than the regime-induced effect (regime  $|d| \approx 0.3$ ). That is, the policy’s trade frequency is determined primarily by the *reward form itself* rather than by market volatility. This is the action-level confirmation of H3 “selective entry  $\rightarrow$  conservative cluster.”

(ii) **Regime adaptation: a consistent decrease across all variants.** Going from low\_vol to high\_vol, all four variants drop trade frequency by about 30% (e.g.,  $\text{sym}$  0.073  $\rightarrow$  0.047). This is mechanically imposed by the ATR-proportional grid formula: as volatility rises, grid widths widen automatically and fill probability falls. The fact that this pattern is the same for all four variants tells us that the cluster separation does *not* originate in differential regime-adaptation capability but rather in differential baseline trade frequency.

### 6.3 Hold Rate: The Mechanism of DSR’s Sliding-Window Memory

The single most striking row in Table 7 is  $\text{dsr}$ ’s hold rate. With trade frequency nearly identical to  $\text{sym}$ ’s (0.073 vs 0.073),  $\text{dsr}$ ’s hold rate of 0.142 is  $1.9\times$   $\text{sym}$ ’s 0.074,  $3.9\times$   $\text{asym}$ ’s 0.036, and  $4.6\times$   $\text{pt}$ ’s 0.031. At high\_vol the gap widens to 0.120 vs 0.020, i.e.,  $6.0\times$   $\text{pt}$ .

The mechanism follows directly from the DSR formulation of §3.4. The DSR step reward

$$r_t^{\text{dsr}} = \frac{B_{t-1} \Delta A_t - \frac{1}{2} A_{t-1} \Delta B_t}{(B_{t-1} - A_{t-1}^2)^{3/2}}$$

is built from the EWMA-window ( $\eta = 1/28$  h,  $\approx 1.2$  days) mean and variance moments. A short holding (buy followed quickly by a sell) injects large noise into both  $\Delta A$  and  $\Delta B$  within the window, and when the denominator  $(B - A^2)^{3/2}$  is small (i.e., variance is small) that noise is amplified in the reward signal. From the policy’s point of view, this reads as “*short holdings have a noisy, weak training signal,*” and the policy preferentially learns *longer holdings*. The other three variants ( $\text{sym}$ ,  $\text{asym}$ ,  $\text{pt}$ ) have step rewards that are instantaneous functions of 1-step PnL with no direct dependence on holding length, hence their low hold rates.

This mechanism aligns with  $\text{dsr}$ ’s top ranking in final Sharpe (at the 1M step) and in cumulative return reported in §5. Long holdings absorb late-stage policy variability and stabilize final performance, and they are the principal driver of  $\text{dsr}$ ’s reversal to 1st place in CPCV (§7.2). At the same time, the very same “preference for long holding” becomes a *fatal liability* in the OOS bull market of §7.3, completing this paper’s central in-sample  $\leftrightarrow$  OOS trade-off causal chain.

### 6.4 Per-Regime Action Distribution

Table 8 and Figure 4 make the across-variant behavioral differences quantitative.

Table 8: Per-regime action statistics by variant.  $\bar{a}^{(0)}$  = aggressiveness,  $\bar{a}^{(1)}$  = profit\_target; mean (and std for  $\bar{a}^{(0)}$ ).

| Variant | Vol regime | $\bar{a}^{(0)}$ mean | $\bar{a}^{(0)}$ std | $\bar{a}^{(1)}$ mean |
|---------|------------|----------------------|---------------------|----------------------|
| sym     | low        | 0.129                | 0.151               | 0.063                |
|         | high       | 0.125                | 0.148               | 0.068                |
| dsr     | low        | 0.129                | 0.217               | 0.145                |
|         | high       | 0.124                | 0.207               | 0.161                |
| asym    | low        | 0.325                | 0.300               | 0.029                |
|         | high       | 0.281                | 0.271               | 0.032                |
| pt      | low        | 0.446                | 0.311               | 0.056                |
|         | high       | 0.399                | 0.301               | 0.053                |

**Buy-side grid position ( $\bar{a}^{(0)}$ ).** sym/dsr settle around  $\bar{a}^{(0)} \approx 0.12$ , placing buy quotes very close to the current price (in §3.1’s grid formula  $g^{\text{buy\_hi}} = r_t(A_b + B_b a^{(0)})$ , small  $a^{(0)}$  gives a narrow buy spacing and a fill-friendly position). asym sits at  $\bar{a}^{(0)} \approx 0.30$  and pt at  $\approx 0.42$ , placing quotes “deeper” away from the current price. asym/pt therefore receive fills only on larger price drops, and their trade frequency falls accordingly (consistent with Table 7).

**Sell-side grid position ( $\bar{a}^{(1)}$ ).** sym, asym, and pt all sit at  $\bar{a}^{(1)} \approx 0.03$ – $0.07$ , with very tight (fast-exit) sell quotes. Only dsr sits at  $\bar{a}^{(1)} \approx 0.15$ , a  $2$ – $3\times$  wider sell spread. This is the direct cause of dsr’s long hold rate (Table 7, right): deeper sell quotes wait longer for a fill and extend holding time.

**Mild regime adaptation.** Going from low\_vol to high\_vol, all four variants reduce  $\bar{a}^{(0)}$  slightly (e.g., pt:  $0.446 \rightarrow 0.399$ ,  $-10\%$ ). Higher volatility nudges policies toward closer buy quotes. However, the change is much smaller than the across-variant differences ( $\approx 0.30$ ), reinforcing the conclusion of §6.2: **the contribution of the vol regime to cluster separation is minor, with the reward form being dominant.**

## 6.5 Counterfactual State-grid Verification

The analysis so far is restricted to the state distribution that policies actually encounter. To complement this, we partition the (atr\_ratio, divergence) plane into a state grid and visualize each variant’s mean action in Figure 5 on the same grid. Observations:

- All four variants behave differently in the region with large divergence ( $s_t^{(1)}$ ) values (deep-loss region relative to average cost), i.e., state-conditional behavioral differences exist beyond just distribution means.
- The cell with the largest across-variant difference is  $s_t^{(1)} < 0$  + moderate atr\_ratio, which corresponds to “the held position is above the average cost, with moderate volatility.” In this region, sym/dsr pick narrow sell grids whereas asym/pt pick hold/wider-sell behavior, i.e., they *wait longer for additional profit*.

This counterfactual analysis confirms directly that the cluster separation arises from differences in the policy mapping  $\pi_{\text{variant}}(a|s)$  itself, not merely from differences in the state distribution

that the policies encounter.

## 6.6 Summary of the Causal Chain

We summarize the chain linking the results of this section to the risk-profile trade-off of §5 in a single line:

Reward form  $\rightarrow$  trade frequency + hold time  $\rightarrow$  Risk profile cluster  $\rightarrow$  (Sharpe, MDD) trade-off

Quantitative evidence per arrow:

1. **Reward form  $\rightarrow$  trade frequency:** The loss-asymmetric variants (**asym**  $\beta = 3.42$ , **pt**  $\lambda = 3.30$ ) push buy quotes deeper ( $\bar{a}^{(0)}$  0.30–0.45 vs  $\sim 0.12$  for **sym/dsr**), yielding 25–35% lower trade frequency (Table 7).
2. **Reward form  $\rightarrow$  hold time:** The DSR sliding-window memory raises the signal-to-noise ratio of longer holdings, so the policy learns a wider sell grid ( $\bar{a}^{(1)} \approx 0.15$ ) and 2–6 $\times$  longer holds. This effect is *independent* of the loss-asymmetry effect; the evidence is that **dsr** matches **sym** in trade frequency but is much higher on hold rate alone.
3. **Trade frequency + hold time  $\rightarrow$  Risk profile:** Few trades + short hold  $\rightarrow$  low MDD + high Calmar (conservative cluster). Frequent trades + short hold  $\rightarrow$  high Sharpe + high MDD (**sym**). Frequent trades + long hold  $\rightarrow$  stable final Sharpe + larger MDD (**dsr**). These three combinations map directly to the per-variant metric leaders in §5.2.
4. **Policy distance:** This behavioral structure is captured statistically by the 2.22 $\times$  policy-distance ratio of §5.5, providing the action-level basis for the distributional cluster separation (within  $|d| < 0.30$ , across  $|d| > 0.79$ ).

Hypothesis **H3** (**selective entry  $\rightarrow$  conservative cluster**) is *strongly supported* by combining the trade-frequency quantification of this section with the buy-grid-position evidence. The section also provides the starting point of this paper’s central in-sample  $\leftrightarrow$  OOS trade-off mechanism (the “DSR’s long holding is both an in-sample asset and an OOS liability” causal chain that returns in §7.3).

## 7 Robustness Analysis

### 7.1 Slippage Robustness

The results of §5 and §6 were obtained on an environment with only a 0.05% fee. As a standard sim2real stress test we now add slippage  $\sigma = 0.02\%$  and ask whether Scenario D (two clusters, Pareto frontier) is preserved. This is the first piece of evidence on hypothesis **H4** (**slippage / CPCV robustness**).

### 7.1.1 Setup

We extend the transaction-cost model of §3.2 with slippage  $\sigma = 0.02\%$ , modifying fill prices to  $q^{\text{buy}} \cdot (1 + \sigma)$  and  $q^{\text{sell}} \cdot (1 - \sigma)$ . This is a worst-case estimate for Binance taker side and corresponds to the 95% quantile of empirically reported slippage on BTC/USDT 1-hour limit fills. All other hyperparameters (formula\_coefs,  $n_{\text{splits}}$ , PPO stabilization package) are *identical* to exp032b in §5. Total: 4 variants  $\times$  10 seeds  $\times$  1M steps = 40 runs, completing in approximately 3 h 47 min.

### 7.1.2 Per-variant Results: Uniform $\sim 12\%$ Decay

Table 9: exp033 results. 10-seed mean  $\pm$  std. The last two columns compare to ATR baseline outperformance (no-slip / with-slip). Phase 16a re-measures the ATR baseline under the same slippage to Val Sharpe 0.835 (MDD 15.27%).

| Variant | exp033 Sharpe<br>(slip 0.02%) | exp032b Sharpe<br>(no slip) | $\Delta$ | Cohen’s $d$<br>(vs exp032b) | Retention | vs ATR-slip<br>(0.835) |
|---------|-------------------------------|-----------------------------|----------|-----------------------------|-----------|------------------------|
| sym     | $1.658 \pm 0.30$              | 1.871                       | −0.213   | −0.80                       | 88.6%     | +99%                   |
| dsr     | $1.551 \pm 0.26$              | 1.809                       | −0.259   | −1.10                       | 85.7%     | +86%                   |
| asym    | $1.478 \pm 0.10$              | 1.681                       | −0.203   | −2.01                       | 87.9%     | +77%                   |
| pt      | $1.459 \pm 0.10$              | 1.667                       | −0.207   | −2.18                       | 87.6%     | +75%                   |

The most striking pattern in Table 9 is that the Sharpe decay is *nearly identical across variants*. The absolute drop falls in the narrow range −0.20 to −0.26, and retentions (exp033/exp032b) cluster in a 3 percentage-point band of 85.7% to 88.6%. The across-variant effect sizes ( $|d| = 0.80$  to 2.18) are all significant in direction (each variant’s drop is statistically reliable across seeds), but the differences *between* variants in their retentions are small. Slippage thus acts as a *cluster-insensitive uniform decay*.

### 7.1.3 MDD Is Almost Unchanged

Table 10: MDD change (exp032b  $\rightarrow$  exp033). *The MDD of the conservative cluster is essentially unchanged.*

| Variant | exp032b MDD (%) | exp033 MDD (%) | $\Delta$ |
|---------|-----------------|----------------|----------|
| sym     | 3.27            | 3.47           | +0.20    |
| dsr     | 4.33            | 4.88           | +0.55    |
| asym    | 2.28            | 2.28           | +0.01    |
| pt      | 2.31            | 2.34           | +0.03    |

asym and pt have  $\Delta < 0.05$  pp in MDD after adding slippage. sym/dsr show a small increase (+0.20, +0.55). This matches the mechanism of §6: slippage is a per-trade cost, so the aggressive cluster (more trades) is more exposed, while the conservative cluster’s lower trade count limits the cumulative exposure.

Table 11: Cluster preservation across robustness tests. Ratios are across/within Cohen’s  $d$  (Sharpe distribution level) or  $L_2$  action distance (policy level).

|                                       | within $ d $ (mean) | across $ d $ (mean) | Ratio        |
|---------------------------------------|---------------------|---------------------|--------------|
| exp032b (no slippage, Sharpe level)   | 0.30                | 0.79                | —            |
| exp033 (slippage 0.02%, Sharpe level) | 0.288               | 0.630               | <b>2.19×</b> |
| exp032c (policy distance $L_2$ )      | 0.129               | 0.286               | <b>2.22×</b> |

#### 7.1.4 Cluster Preservation: Scenario D Is Robust

Table 11 is the central result of this section. After adding slippage, the across/within Cohen’s  $d$  ratio is  $2.19\times$ —essentially identical to the  $2.22\times$  ratio of §5.5 (policy distance) and consistent with exp032b. Therefore **the two-cluster structure of Scenario D is robust to slippage**. The “risk-profile-dimension trade-off” identified in §5 does not depend on the precise transaction-cost model.

#### 7.1.5 Fair Comparison Against the ATR Baseline (Phase 16a)

The initial analysis kept the ATR baseline at its no-slippage value 1.378 and evaluated RL under slippage, an unfair comparison. In this setup the conservative cluster (**asym** 1.478, **pt** 1.459) only marginally exceeds ATR 1.378, and a caveat was needed. Phase 16a (May 16, 2026) re-measured the ATR baseline under the same slippage (0.02%), giving **Val Sharpe** 0.835 (**MDD** 15.27%), a roughly  $-39\%$  drop. The reason is that the ATR baseline policy uses both buy and sell quotes at every step and therefore incurs more cumulative slippage than the RL policies.

A fair comparison against ATR-with-slippage 0.835 gives:

- **sym** 1.658:  $+99\%$  vs **ATR**
- **dsr** 1.551:  $+86\%$  vs **ATR**
- **asym** 1.478:  $+77\%$  vs **ATR**
- **pt** 1.459:  $+75\%$  vs **ATR**

Under slippage, **all four RL reward variants exceed the ATR baseline by 75–99%**, and the initial caveat “the conservative cluster does not reach ATR” disappears under fair comparison. **H2 weak (asym/dsr/pt > ATR)** is *strongly supported* under slippage.

#### 7.1.6 Section Summary and Partial Verdict on H4

- Slippage 0.02% causes a uniform  $\sim 12\%$  Sharpe decay across all four reward variants. The two-cluster structure is preserved quantitatively (across/within ratio  $2.19\times \approx 2.22\times$ ).
- Conservative-cluster MDD is essentially unchanged; aggressive-cluster MDD rises slightly. The trade-frequency gap of §6 carries over to differential slippage exposure.
- Re-evaluating ATR under the same slippage drops its Sharpe to 0.835, and all four RL variants now exceed ATR by 75–99%. **H2 weak is strongly supported**.
- **H4 (slippage robustness)**: *partially supported*. The qualitative cluster-preservation finding

( $2.19\times$ ) is robust; absolute Sharpe values undergo uniform decay.

The next subsection (§7.2) asks the same robustness question in the multi-split setting (6-fold CPCV, 15 paths).

## 7.2 Multi-Split Evaluation via CPCV

[TBD. *exp034: 6-fold CPCV 15 paths. DSR  $p < 0.001$  (Bonferroni 4-way correction);  $\mathbf{dsr}$  reversal to 1<sup>st</sup> (1.413, 5% CVaR 0.890); cluster preservation ratio  $3.55\times$ .]*

## 7.3 Out-of-Sample Evaluation and Prospect Theory’s Robustness

[TBD. *Central chapter. exp035: 100 RL models + ATR baseline + Buy-and-Hold on Test.  $\mathbf{pt}$  1<sup>st</sup> on Test (both sources  $p < 0.002$ );  $\mathbf{dsr}$  last on Test. Phase 16d mechanism: hold-duration distribution (DSR 7-day max vs.  $\mathbf{pt}$  6h max); causal account of  $\mathbf{pt}$ ’s OOS robustness.]*

## 8 Discussion

[TBD. *Subsections: §8.1 Val–Test distribution-shift quantification ( $KS\ p < 10^{-10}$ ,  $-27\%$  variance); §8.2 honest acknowledgment of environment dependence in *exp027\_rl* prior evidence (*Env-v3 asym* 1.955); §8.3 limits of single-metric winner; §8.4 limitations: BTC-only, single regime in Test, single slippage level.]*

## 9 Conclusion

[TBD. *Four main messages: (1) Pareto-frontier trade-off; (2) winner reversal across evaluation environments; (3) OOS robustness of Prospect Theory; (4) in-sample advantage  $\neq$  OOS robustness. Future work: multi-asset, additional regimes, DR, meta-RL.]*

## References

- Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. Optuna: A next-generation hyperparameter optimization framework. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 2623–2631, 2019.
- Marco Avellaneda and Sasha Stoikov. High-frequency trading in a limit order book. *Quantitative Finance*, 8(3):217–224, 2008. doi: 10.1080/14697680701381228.
- David H. Bailey and Marcos López de Prado. The deflated sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5):94–107, 2014.
- Charan Bandaruwalli. Risk-aware deep reinforcement learning for cryptocurrency and equity trading. *arXiv preprint*, 2025.

- Bryan Gort, Xiao Chen, Matthew Tong, Shaojie Xu, and Xingjian Li. Deep reinforcement learning for cryptocurrency trading: Practical approach to address backtest overfitting. *arXiv preprint arXiv:2209.05559*, 2022.
- Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. Deep reinforcement learning that matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32, 2018.
- Daniel Kahneman and Amos Tversky. Prospect theory: An analysis of decision under risk. *Econometrica*, 47(2):263–291, 1979. doi: 10.2307/1914185.
- Hao Yuan Timothy Liu. Reinforcement learning trading strategies: A literature review. *arXiv preprint arXiv:2502.06551*, 2025.
- Yang Liu, Qi Liu, Hongke Zhao, Zhenya Pan, and Chuanren Liu. Adaptive quantitative trading: An imitative deep reinforcement learning approach. 34(02):2128–2135, 2020. Referenced as Liu et al. (2021) for the PPO+LSTM Bitcoin trading variant.
- Marcos López de Prado. *Advances in Financial Machine Learning*. Wiley, 2018. Combinatorial Purged Cross-Validation (CPCV), Chapter 7.
- John Moody and Matthew Saffell. Learning to trade via direct reinforcement. *IEEE Transactions on Neural Networks*, 12(4):875–889, 2001. doi: 10.1109/72.935097. Differential Sharpe Ratio (DSR) formulation.
- Andrew Y. Ng, Daishi Harada, and Stuart Russell. Policy invariance under reward transformations: Theory and application to reward shaping. In *Proceedings of the 16th International Conference on Machine Learning (ICML)*, pages 278–287, 1999.
- Minh Pham, Anh Le, and Binh Nguyen. Dqn-based bitcoin technical trading strategy selection. *arXiv preprint*, 2025.
- Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus, and Noah Dormann. Stable-baselines3: Reliable reinforcement learning implementations. *Journal of Machine Learning Research*, 22(268):1–8, 2021.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Siyu Sun, Rundong Wang, and Bo An. Reinforcement learning for quantitative trading. *ACM Transactions on Intelligent Systems and Technology*, 14(3):1–29, 2023. doi: 10.1145/3582560.
- Amos Tversky and Daniel Kahneman. Advances in prospect theory: Cumulative representation of uncertainty. *Journal of Risk and Uncertainty*, 5(4):297–323, 1992.  $\alpha = 0.88$ ,  $\lambda = 2.25$  standard parameters.
- J. Welles Wilder. New concepts in technical trading systems. 1978.
- Hafiz Muhammad Athar Yasin and Asim Anwar Gill. Reinforcement learning framework for quantitative trading. *International Journal of Computer Applications*, 186(8):1–9, 2024.

Zihao Zhang, Stefan Zohren, and Stephen Roberts. Deep reinforcement learning for trading.  
*The Journal of Financial Data Science*, 2(2):25–40, 2020.

exp032c Menu 2 — Action distribution per variant x volatility regime  
(red star = mean. n = 10 seeds x ~8700 steps per cell)




Figure 4: 2D action distribution ( $a^{(0)}, a^{(1)}$ ) over the 4 variants  $\times$  3 vol regimes grid. Points are trajectory steps; red crosses are per-cell means. **sym/dsr** cluster around  $a^{(0)} \approx 0.12$  (narrow buy spacing); **asym/pt** cluster around  $a^{(0)} \approx 0.30\text{--}0.45$  (deeper buy quotes). Only **dsr** has  $a^{(1)} \approx 0.15$ , with a slightly wider sell grid to extend holding time.




Figure 5: Counterfactual mean-action heatmaps over the (atr\_ratio, divergence) state grid for each variant. All four variants differ most in the region with large negative divergence and mid atr\_ratio (i.e., when the held position is above the average cost and volatility is moderate).




Figure 6: Side-by-side comparison of exp032b (light) vs exp033 (dark). Mean  $\pm$  std bars are overlaid with the 10 seed points for each variant. All four variants drop mean Sharpe by roughly 0.20 under slippage, but the ordering and the cluster separation are preserved.